


HYPR.AI

Cloud-based AutoML


Ryan Reece
Insight AI Fellow

Also with Jiangming “Jimmy” Hu
Previous Insight Fellow



Intro

- AI is an *experimental* science
- Search & test many models: “*hyperparameters*”
- Better tools can streamline this process
- Leads to discovering and organizing better models
- Jeff Dean: tools for “ML 2.0” or “*AutoML*”




Google Cloud Platform


H₂O.ai

Typical ML workflow today

Models



Compute



```
Train the model!
This is the fun part: you can feed the training data loaded in earlier into this model and it will learn to classify digits

In [9]: model.fit(X_train, Y_train,
                  batch_size=128, nb_epoch=4,
                  show_accuracy=True, verbose=1,
                  validation_data=(X_test, Y_test))

Train on 60000 samples, validate on 10000 samples
Epoch 1/4
60000/60000 [=====] - 10s - loss: 0.2521 - acc: 0.9245 - val_loss: 0.1131
- val_acc: 0.9651
Epoch 2/4
60000/60000 [=====] - 10s - loss: 0.1016 - acc: 0.9687 - val_loss: 0.0827
- val_acc: 0.9746
Epoch 3/4
60000/60000 [=====] - 11s - loss: 0.0711 - acc: 0.9779 - val_loss: 0.0668
- val_acc: 0.9791
Epoch 4/4
60000/60000 [=====] - 11s - loss: 0.0557 - acc: 0.9816 - val_loss: 0.0642
- val_acc: 0.9805
Out[9]: <keras.callbacks.History at 0x12655bbe0>

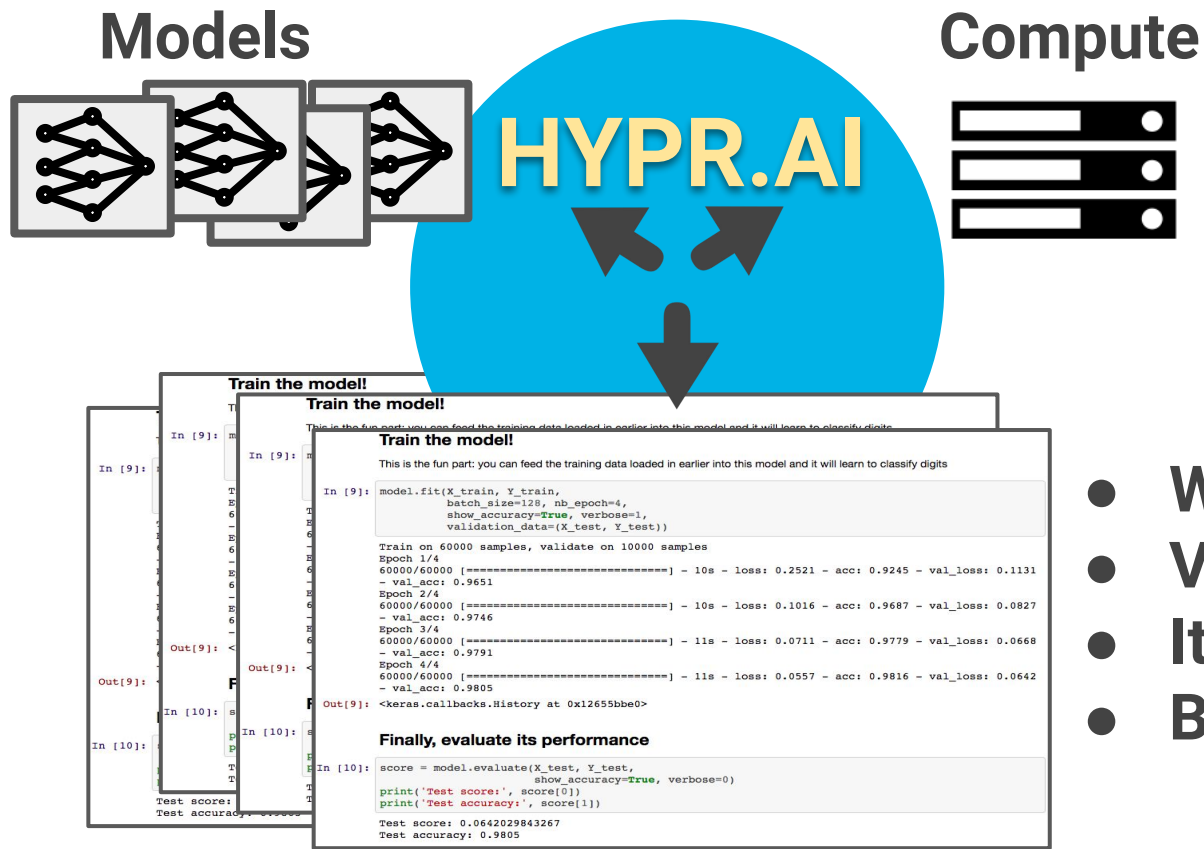
Finally, evaluate its performance

In [10]: score = model.evaluate(X_test, Y_test,
                                show_accuracy=True, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

Test score: 0.0642029843267
Test accuracy: 0.9805
```

- Workflow
- Visualization
- Iterate
- Bookkeeping?

Typical ML workflow today



- Workflow
- Visualization
- Iterate
- Bookkeeping?


HYPR.AI provides




Website UI: construct or upload keras models




Database for storing models and results



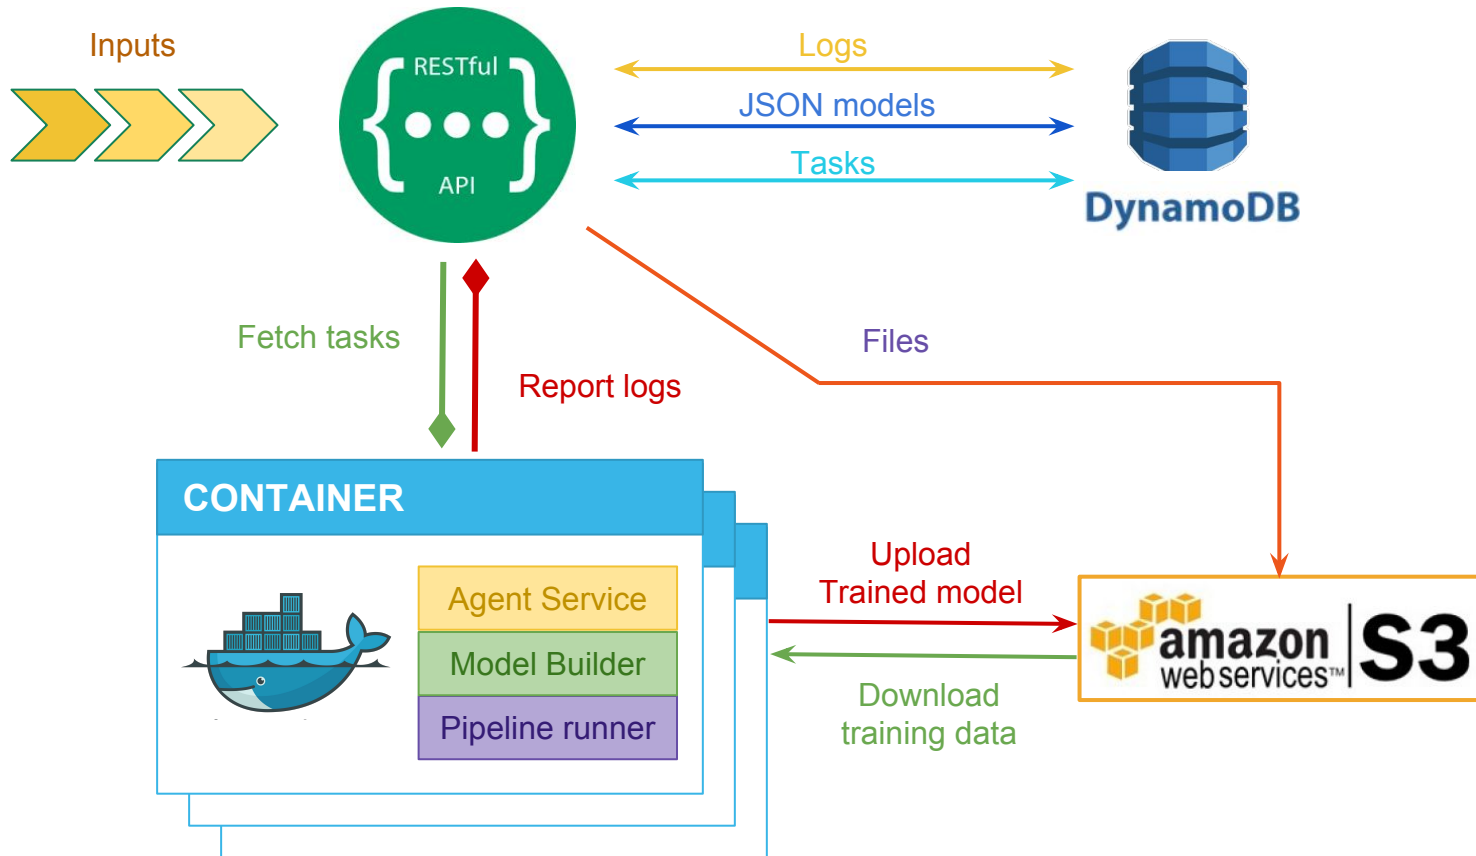
Task queue for training many models



Compute back-ends for AWS & Paperspace



You





HYPR.AI

Cloud-based AutoML

Models

Tasks

Results

Storage

Cluster

Settings



Models

+ New



CNN-Base

[Convolution2D](#) (conv1) - inputs: [null,32,32,3] filte...

[Convolution2D](#) (conv2) - filters: 64 kernel_size: [3,...

[MaxPooling2D](#) (pool1) - pool_size: [2,2] strides: [2,...

[Dropout](#) (dropout1) - rate: 0.25

[Convolution2D](#) (conv3) - filters: 128 kernel_size: [3,...

[MaxPooling2D](#) (pool2) - pool_size: [2,2] strides: [2,...

[Convolution2D](#) (conv4) - filters: 128 kernel_size: [3,...

[MaxPooling2D](#) (pool3) - pool_size: [2,2] strides: [2,...

[Dropout](#) (dropout2) - rate: 0.25

[Flatten](#)

[Dense](#) (dense1) - units: 1024 activation: relu

[Dropout](#) (dropout3) - rate: 0.5

[Dense](#) (softmax1) - units: 10 activation: softmax

[Compiler](#) optimizer: {"adam":{"lr":0.0001}} loss: cat...



CNN-Base-100

[From](#): CNN-Base

[Dense](#) (softmax1) - units: 100 activation: softmax

[Compiler](#) optimizer: {"adam":{"lr":0.0001}} loss: cat...



HYPR.AI

Cloud-based AutoML

+ Model

Task

Re

Storage

Cloud



Models

+ New



CNN-Base

Convolution2D (conv1) - inputs: [null,32,32,3] filter...


Convolution2D (conv2) - filters: 64 kernel_size: [3,...



CNN-Base-100

From: CNN-Base

Dense (softmax1) - units: 100 activation: softmax



Dropout (dropout3) - rate: 0.5

Dense (softmax1) - units: 10 activation: softmax

Compiler optimizer: {"adam":{"lr":0.0001}} loss: cat...



HYPR.AI

Cloud-based AutoML

Models

Tasks

Results

Storage

Cluster

Settings

Add new hyper task



Task Name

my-test

Max Epoch

60

JSON model

CNN-Base

Dataset

cifar-10

Hyperparameters

{'conv2/filters' : ['ci', [32, 64, 128]], 'dropout2/rate' : ['cf', [0.25, 0.50]]}

Weights (pretrained model)

NONE

Cancel

Save

Created

2018-02-08 02:33

2018-02-08 02:23

2018-02-07 03:42

2018-02-07 03:33

2018-02-07 03:29

2018-02-07 01:58

2018-02-07 01:30



test7

CNN-Base

cifar-10

20

completed



HYPR.AI

Cloud-based AutoML

Models

Tasks

Results

Storage

Cluster

Settings

	my-test-014	CNN-Base-my-test-014	cifar-10	60	training	2018-02-08 03:21
	my-test-013	CNN-Base-my-test-013	cifar-10	60	initial	2018-02-08 03:21
	my-test-012	CNN-Base-my-test-012	cifar-10	60	initial	2018-02-08 03:21
	my-test-011	CNN-Base-my-test-011	cifar-10	60	initial	2018-02-08 03:21
	my-test-010	CNN-Base-my-test-010	cifar-10	60	initial	2018-02-08 03:21
	my-test-009	CNN-Base-my-test-009	cifar-10	60	initial	2018-02-08 03:21
	my-test-008	CNN-Base-my-test-008	cifar-10	60	initial	2018-02-08 03:21
	my-test-007	CNN-Base-my-test-007	cifar-10	60	initial	2018-02-08 03:21
	my-test-006	CNN-Base-my-test-006	cifar-10	60	completed	2018-02-08 03:21



HYPR.AI

Cloud-based AutoML

Models

Tasks

Results

Storage

Cluster

Settings



Results



Name	Loss (train/test)	Accuracy (train/test)	Best Epoch	Created
my-test-025	0.60749 / 0.64404	0.78616 / 0.7789	43 / 46	2018-02-08 03:21
my-test-024	0.56335 / 0.62973	0.80176 / 0.784	60 / 60	2018-02-08 03:21
my-test-023	0.69146 / 0.66203	0.75706 / 0.7713	60 / 60	2018-02-08 03:21
my-test-022	0.598 / 0.62903	0.78904 / 0.7858	59 / 60	2018-02-08 03:21
my-test-021	0.66935 / 0.65053	0.76508 / 0.7738	59 / 60	2018-02-08 03:21
my-test-020	0.47215 / 0.61166	0.83378 / 0.7902	48 / 49	2018-02-08 03:21
my-test-019	0.75925 / 0.72927	0.73242 / 0.7433	45 / 48	2018-02-08 03:21



HYPR.AI

Cloud-based AutoML

Models

Tasks

Results

Storage

Cluster

Settings



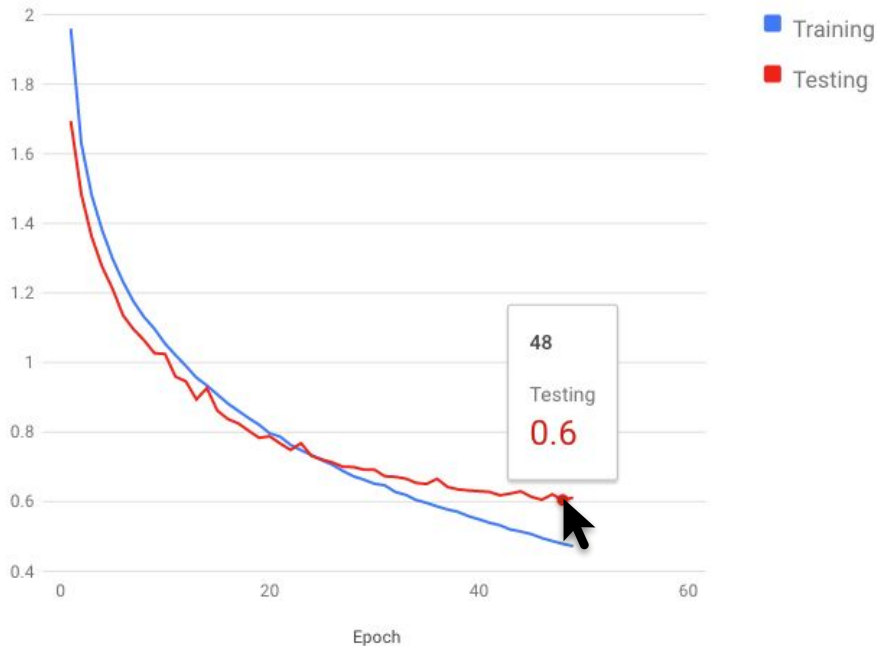
my-test-020

Use model: CNN-Base-my-test-020 ----- Max Epochs: 60

Dataset: cifar-10

Status: completed

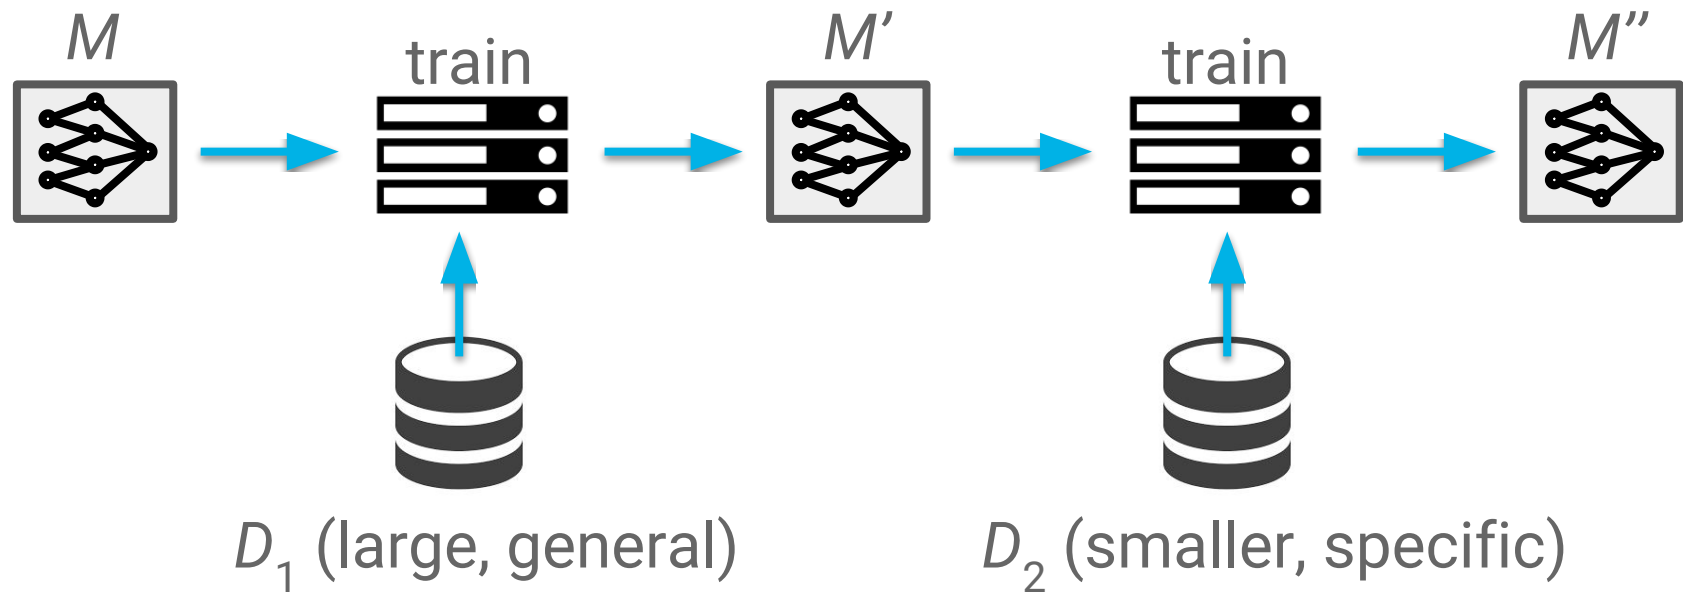
Loss vs epoch



est epoch	Created
3 / 46	2018-02-08 03:21
0 / 60	2018-02-08 03:21
0 / 60	2018-02-08 03:21
9 / 60	2018-02-08 03:21
9 / 60	2018-02-08 03:21
8 / 49	2018-02-08 03:21
5 / 48	2018-02-08 03:21
7 / 50	2018-02-08 03:21
5 / 48	2018-02-08 03:21

Transfer learning

Train a model on one dataset, and then more on another.




Advantages of using HYPR.AI


- Bookkeeping datasets/models/weights in the cloud
- RESTful API allows back-end independence
- Modular models facilitate transfer learning
- Hyperparameter optimization: currently random search or SCaN. Possible future plugins: e.g. Bayesian (SigOpt), Hyperopt
- Try out the running example: <http://hypr.umx.io>

Ryan Reece, Ph.D.
Data scientist, AI/ML/Stats

**Previously a particle physicist
with the ATLAS experiment at
the Large Hadron Collider**



CERN




2006: first summer at CERN

TRT end-cap

electronics


statistics



climbing



Backup slides



Random Search for Hyper-Parameter Optimization

James Bergstra

Yoshua Bengio

Département d'Informatique et de recherche opérationnelle


Université de Montréal

Montréal, QC, H3C 3J7, Canada

JAMES.BERGSTRA@UMONTREAL.CA


YOSHUA.BENGIO@UMONTREAL.CA

Grid Layout




Important parameter

Random Layout



Important parameter



After 100 iterations of my random search alg:


```
best_hparams :  
x = 0.472225  
y = 0.791885  
best_loss = 0.00608015
```




HYPR.AI

Cloud-based AutoML

 Models

 Tasks

 Results

 Storage

Dataset

Weights

 Cluster

Cluster (workers)




Name	Status	Last seen	Server Info
paperspace	idle	2018-01-26 17:36:27	{}
_paperspace-psrxvede4	offline	2018-01-22 21:15:56	{}
_paperspace_psrxvede4	offline	2018-01-22 21:06:16	{}
54.153.100.55	offline	2018-01-19 04:22:31	{}
AWS-g2.2xlarge	offline	2018-01-12 20:18:01	{}
PAPERSPACE-P5000-1	offline	2018-01-12 20:17:58	{}

Data and models

- As proofs of concept, using standard image classification datasets: CIFAR-10, MNIST

1 1 5 4 3
7 5 3 5 3
5 5 9 0 6
3 5 2 0 0

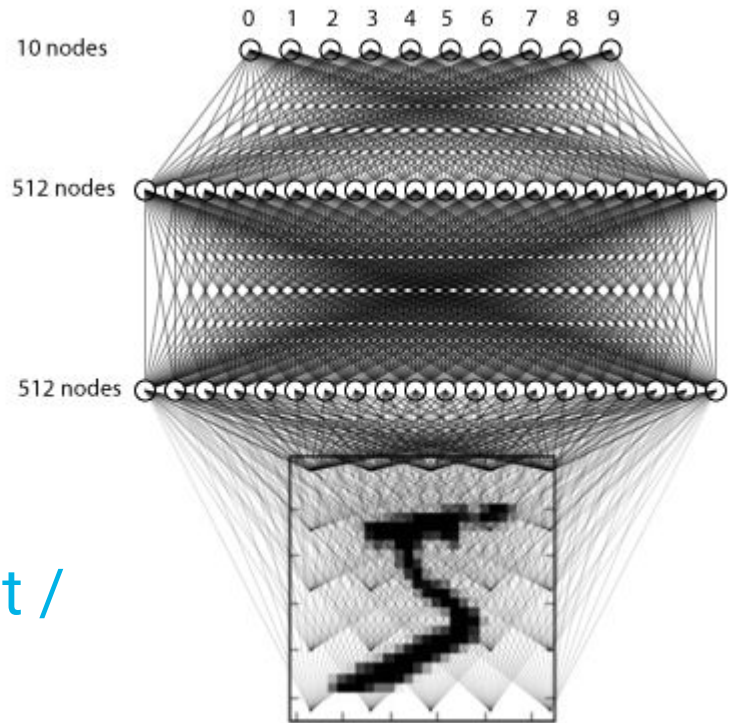
- Successful models use several layers of CNNs + pooling, e.g. LeNet



- With this platform, easily upload and test many architectures

Goals

- Demonstrate a successful hyperparameter scan
- Use our system to discover a performant model
- Future upgrades could add better optimization (e.g. SigOpt / Hyperopt / custom RL?)



LHC/CERN

ATLAS



3-level trigger

40 MHz \rightarrow 100 kHz

\rightarrow 6 kHz \rightarrow 500 Hz

Trigger
& DAQ

raw data

100101011

~ 10 PB/year

ATLAS Data Flow

Worldwide LHC Computing Grid

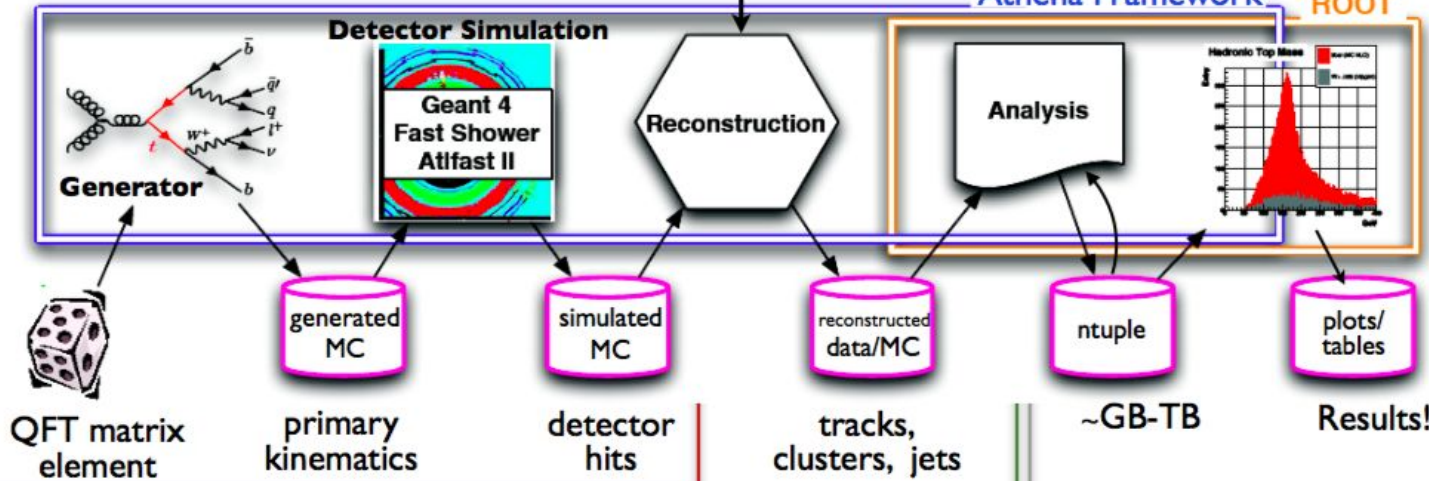
Monte Carlo production

~ 100 k CPUs
over 100 PB


Local resources

Athena Framework

ROOT



2006




Did my first “shift” as we saw the first
cosmics in the few ϕ -slices that were
fully instrumented.


TRT threshold calibration

2008


Before



noise
occupancy



After



- Developed a GUI making it easy for shifters to archive scans to a database for monitoring long-term detector health.
- Still used in the regularly scheduled calibration periods between beam fills.
- Supported TRT as part of DAQ on-call team.

